

redis

a
data structure
server

Rob Tandy
Lunch and Learn
2012-10-16

linked list



element access: scales with # elements $O(n)$

push, pop: fast. independant of n $O(1)$

binary search
tree

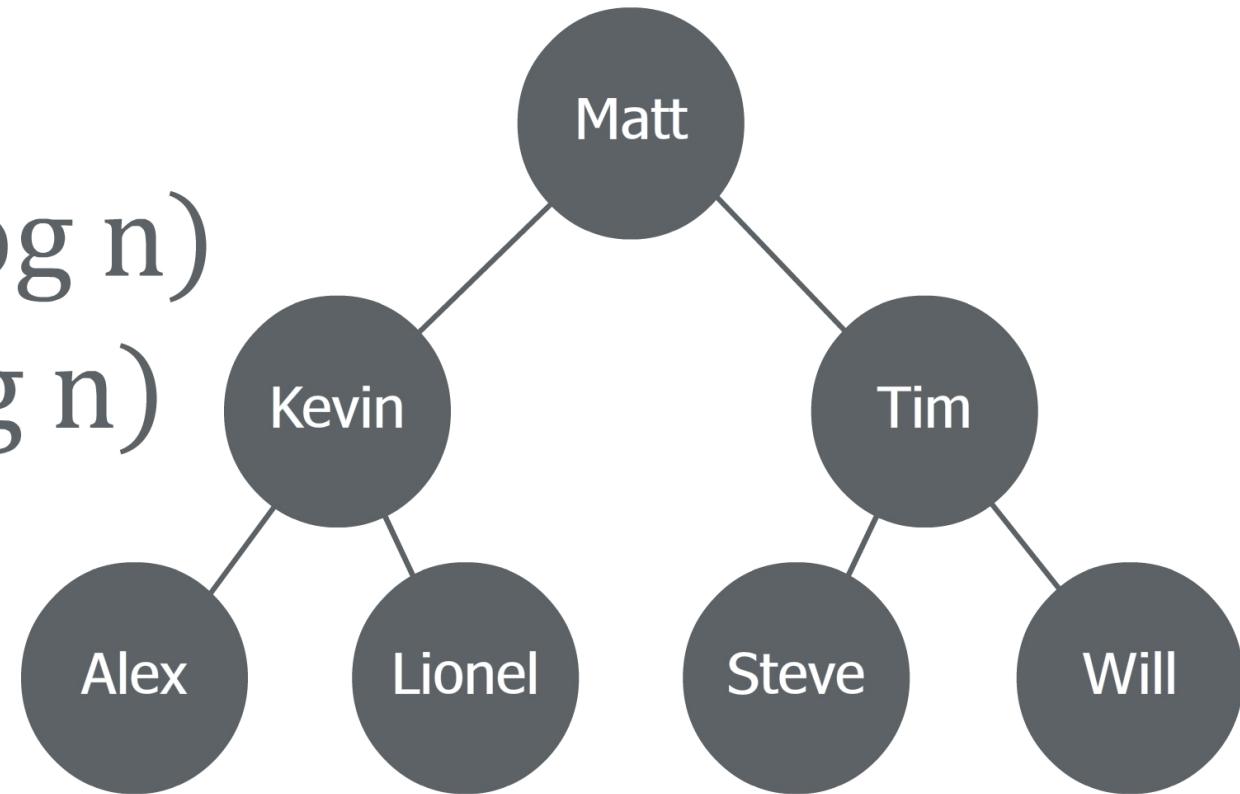
height:

$$2^h = n$$

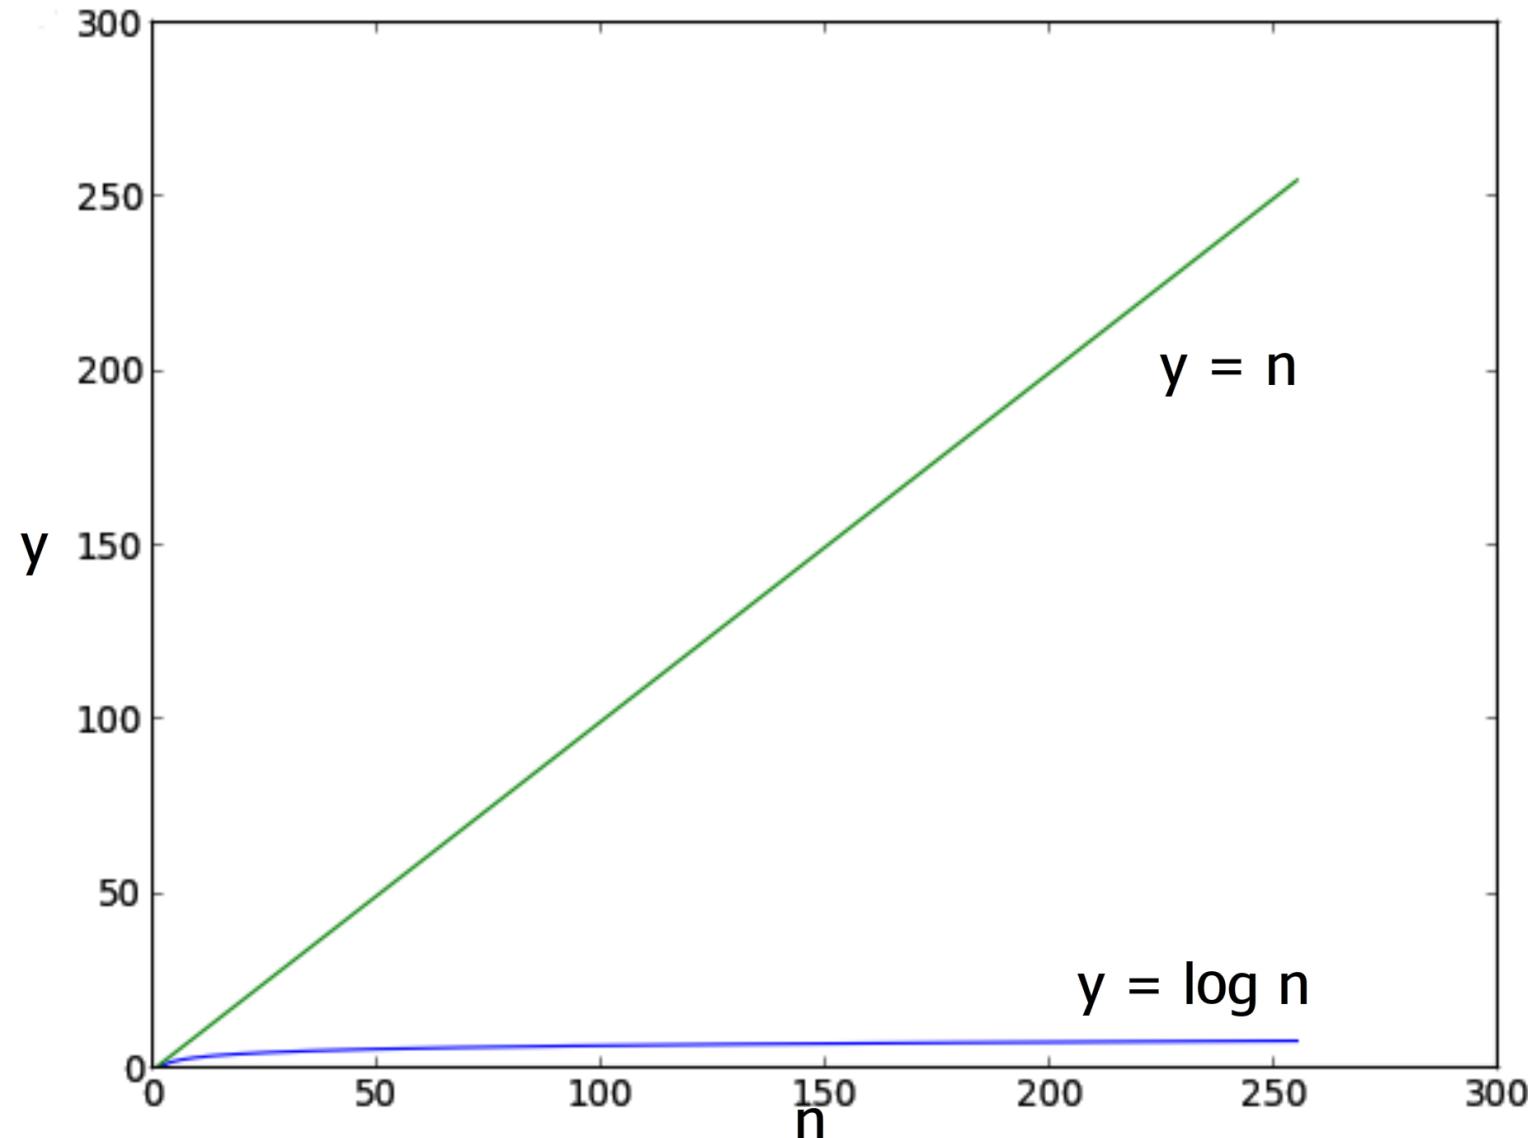
$$h = \log n$$

search: $O(\log n)$

insert: $O(\log n)$



HUGE difference in speed!



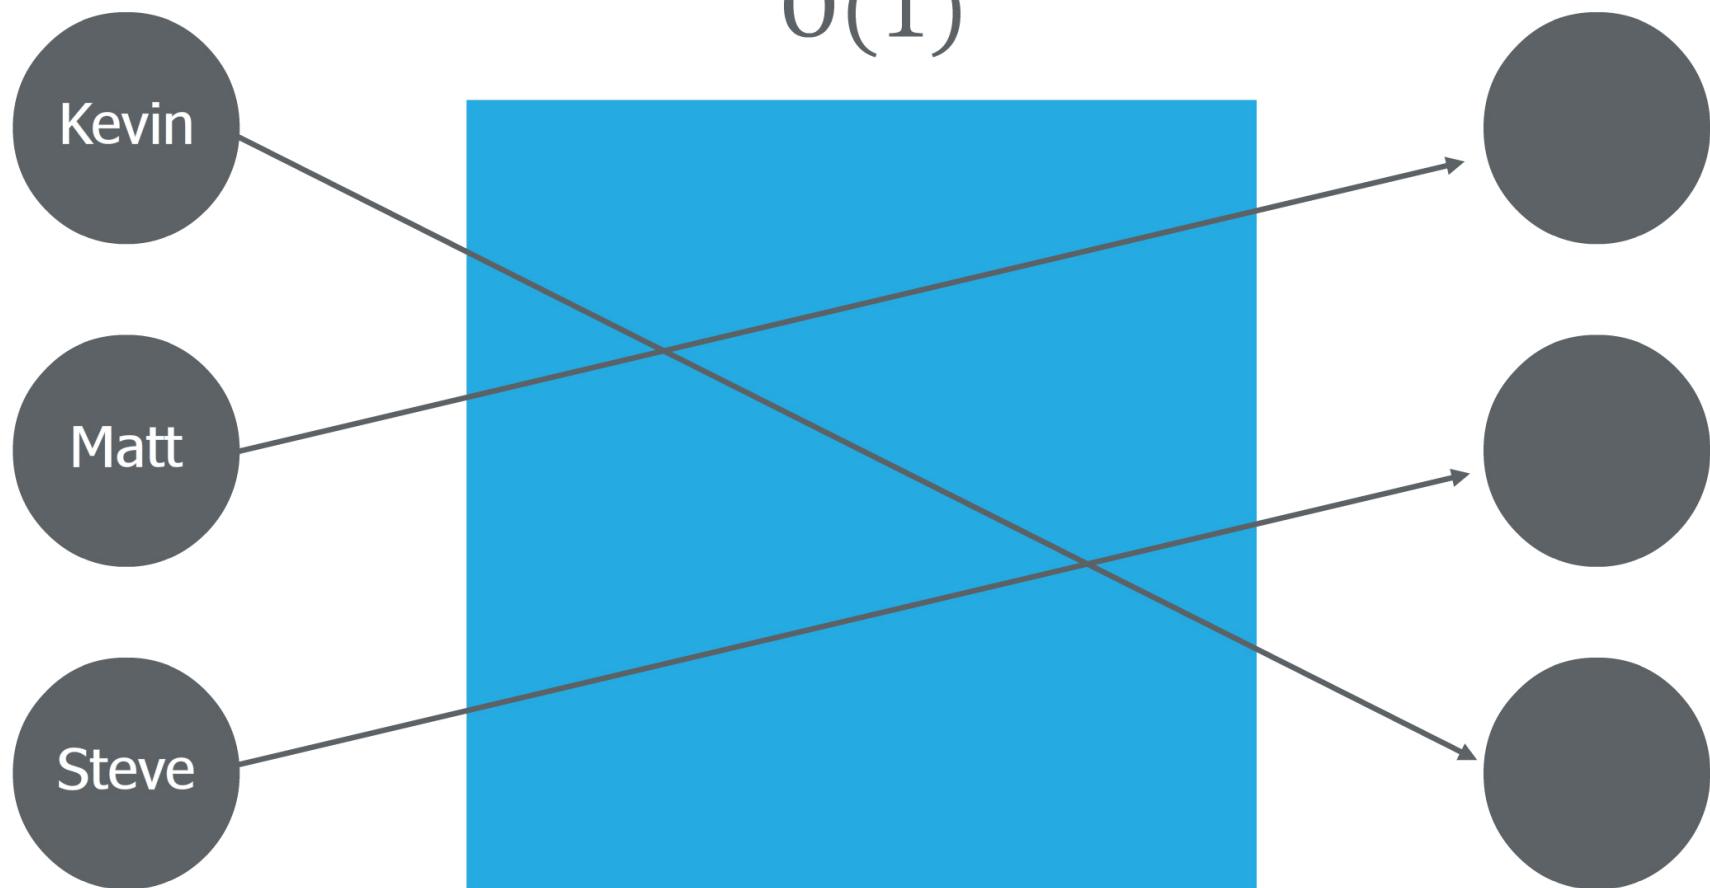
$$\log(33554432) = 25$$

hash function

hash table

Search not dependent on n!

$O(1)$



Like a hash table table only **just keys** instead of
key, value pairs

Still $O(1)$ for access, delete

Quiz: traverse a list and remove duplicates...

```
def remove_dupes(a_list):
    cleaned = []
    seen = set ()
    for item in a_list:
        if not item in seen:
            seen.add (item)
            cleaned.append (item)
    cleaned
```

if $n = 1e9$, CPU can do $1e9$ ops/sec ==> ~ 1 sec. vs ~ 31 years!*

hash tables for everything?

only if you want **FAST** access to your data:

- * caches,
- * associative arrays,
- * implementing objects
(python, ruby),
- * implementing sets

HASH TABLES ARE SO GREAT, I WANT TO MAKE THEM AVAILABLE AS A NETWORK SERVICE!

April 2009: redis (www.redis.io)

- * written by Salvatore Sanfilippo
- * a key-value store
- * simplest working solution for a networked hash table
- * fast. 150k GET/SET per second on i7 laptop
- * March 2010: VMware sponsored

redis is **simple**

Command line interface

```
shell> redis-cli  
redis 127.0.0.1:6379> SET lunch good  
"OK"  
redis 127.0.0.1:6379> GET lunch  
"good"
```

Python - all commands are methods

```
import redis  
  
r = redis.StrictRedis (host='localhost')  
r.set ('lunch' , 'good')  
print r.get ('lunch')  # prints "good"
```

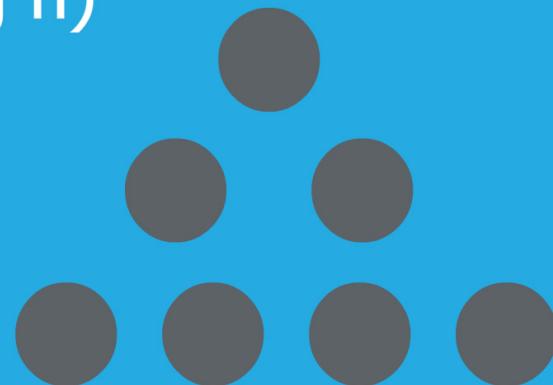
where does it fit?

Database cache

without an index your database table behaves like a list (more or less)



with an index your database table responds like a binary search tree ($\log n$)



`redis:`

- * **faster** and more consistent than an index
- * can cache queries that it doesn't make sense to index
- * more flexible

Your Stuff is safe

- * can log every operation to append-only file - fast
- * fsync at every second (default) or every operation
- * trivially simple master - slave replication

Its a data structure server!

sets:

SADD, SREM, SISMEMBER

lists:

LPOP, LPUSH, RPOP, RPUSH,
LLEN,

BLPOP, BRPOP

Shut up and take my money!

- * sorted sets!
- * publish - subscribe!
- * set expiration for keys!
- * increment / decrement!
- * hashes (one key, many values)!

now where does it fit ?

when you need **FAST**, distributed access to data:

- * web session management (hash: one key: many values)
- * application cache
- * as an internal service of transient data (eg. rate limiter)

when you want to leverage its **simple** treatment of useful design patterns

- * key - value store
- * publish subscribe
- * message broker

Worked Example: A Message Queue

```
def consumer ():

    r = redis.StrictRedis ()

    while True:

        qname, msg = r.blpop ('hello_queue')
        print 'received', msg


def producer (name):

    r = redis.StrictRedis (),
    i = 0

    while True:

        m = 'hi from {0}! msg #{1}'.format (name, i)
        i += 1

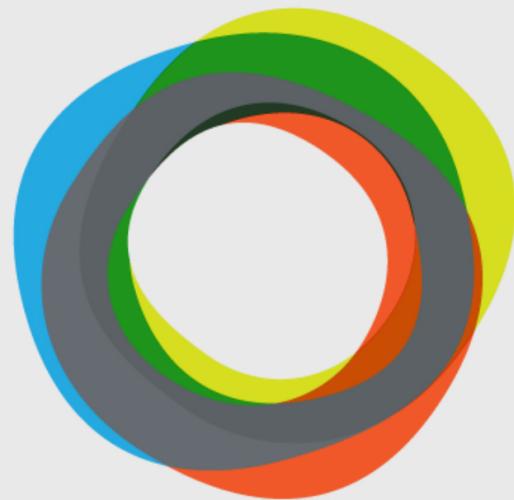
        r.rpush ('hello queue', m)
        time.sleep (5)
```

try more than one producer and consumer!

Worked Example 2 : Pub - Sub

```
def subscriber (channel):  
    pubsub = redis.StrictRedis ().pubsub ()  
    pubsub.subscribe (channel)  
  
    for msg in pubsub.listen ():  
        if msg['type'] == 'message':  
            print 'received' , msg['data'] , 'on' , msg['channel']
```

```
def publisher (name, channel):  
    r = redis.StrictRedis (),  
    i = 0  
  
    while True:  
        m = 'hi from {0}! msg #{1}'.format (name, i)  
        i += 1  
  
        r.publish (channel, m)  
        time.sleep (5)
```



FINDAWAY WORLD